

Lab4-Stat131A-ColinAsbill

AUTHOR


Colin Asbill

```
library(tidyverse)
library(broom)
library(GGally)
library(leaps)
diamonds <- read_csv("diamonds.csv")
```

Part 1 EDA and First Fit

Question 1

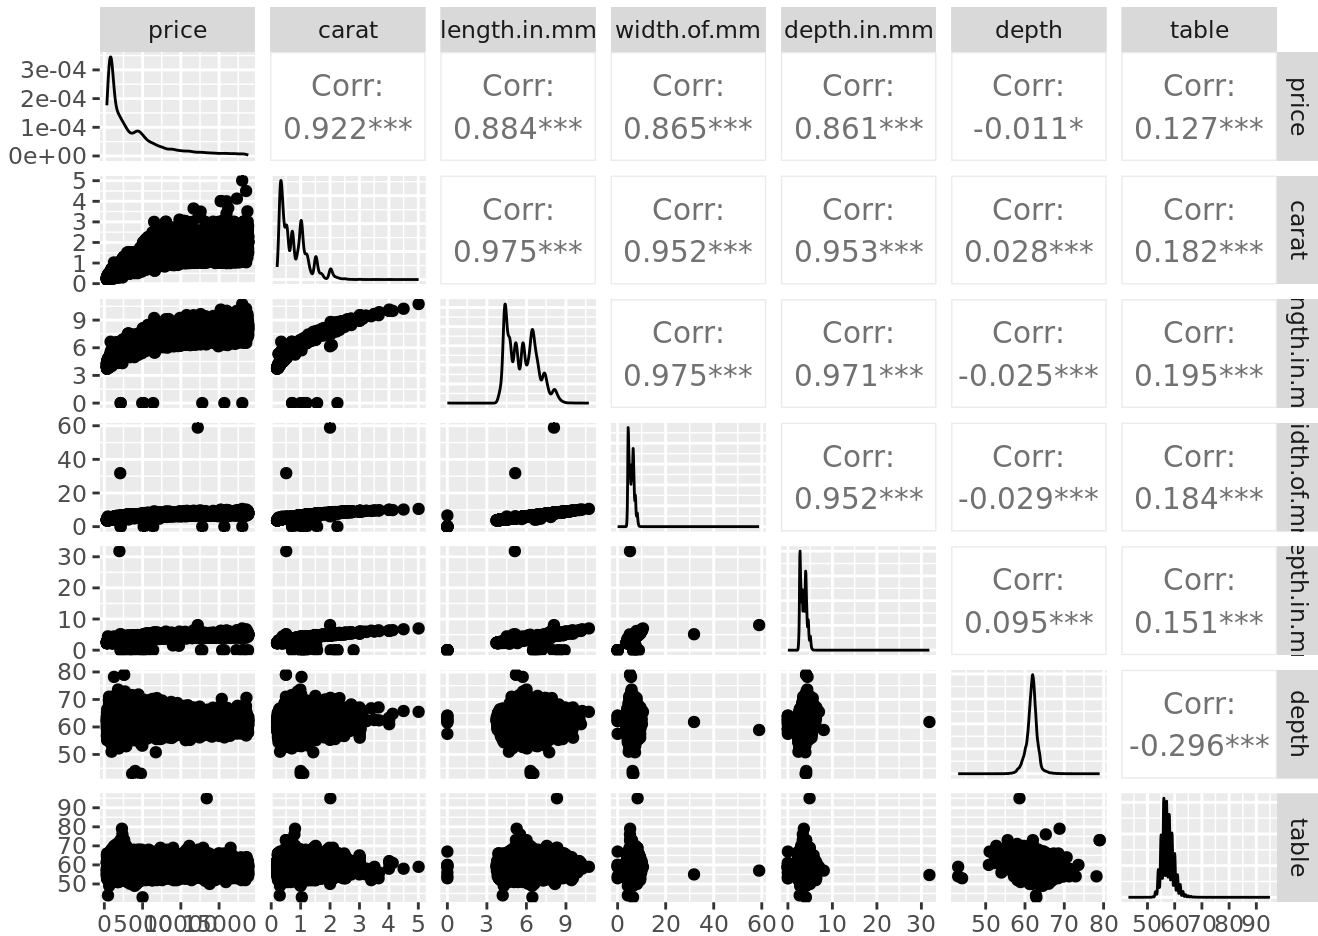
```
diamonds %>%
  ggplot(aes(x = price, y = carat)) +
  geom_point(aes(color = cut))
```




Looking at this graph that compares price with the carat of the diamonds we can see that carat and price have correlation, but the cut variable might be affected by outliers such as the 5 carat diamonds in the top

right being very expensive, but having the lowest quality cut.

```
diamonds %>%
  select(price, carat, length.in.mm, width.of.mm, depth.in.mm, depth, table) %>%
  ggpairs()
```



```
diamonds_new <- diamonds %>% select(price, carat, length.in.mm, width.of.mm, depth.in.mm, depth,
as.data.frame(cor(diamonds_new)) %>%
  rownames_to_column("Variables_1") %>%
  pivot_longer(-c(Variables_1), names_to = "Variables_2", values_to = "Correlation") %>%
  ggplot(mapping = aes(x = Variables_1, y = Variables_2)) +
  geom_tile(aes(fill = Correlation)) +
  scale_fill_gradient(low = "black", high = "red")+
  theme_classic() +
  theme(axis.title.x = element_blank(),
        axis.title.y = element_blank())
```



We can see Price is not very correlated with table and depth, while the rest of the continuous variables have similar levels of correlation

Question 2

```
lm_d <- lm(price ~ carat + depth + table + length.in.mm + width.of.mm + depth.in.mm, data = diamonds)
tidy(lm_d) %>%
  select(term, estimate, std.error)
```

A tibble: 7 × 3

term	estimate	std.error
<chr>	<dbl>	<dbl>
1 (Intercept)	20849.	448.
2 carat	10686.	63.2
3 depth	-203.	5.50
4 table	-102.	3.08
5 length.in.mm	-1316.	43.1
6 width.of.mm	66.3	25.5
7 depth.in.mm	41.6	44.3

Coefficient std.error

```
carat 10686.3091 63.200807
depth -203.1541 5.503984
table -102.4457 3.084213
length.in.mm -1315.6678 43.070264
width.of.mm 66.3216 25.523021
depth.in.mm 41.6277 44.304632
```

Question 3

price = 20849.3164 + 10686.3091carat + -203.1541depth + -102.4457table + -1315.6678length.in.mm + 66.3216width.of.mm + 41.6277depth.in.mm

Coefficient Interpretation

Question 1

For every one unit increase in carat there is a corresponding \$10,686.3091 increase in price

Question 2

```
diamonds$carat_standardized <- (diamonds$carat - mean(diamonds$carat))/(sd(diamonds$carat))
diamonds$depth_standardized <- (diamonds$depth - mean(diamonds$depth))/(sd(diamonds$depth))
diamonds$table_standardized <- (diamonds$table - mean(diamonds$table))/(sd(diamonds$table))
diamonds$length_standardized <- (diamonds$length.in.mm - mean(diamonds$length.in.mm))/(sd(diamonds$length.in.mm))
diamonds$width_standardized <- (diamonds$width.of.mm - mean(diamonds$width.of.mm))/(sd(diamonds$width.of.mm))
diamonds$depthmm_standardized <- (diamonds$depth.in.mm - mean(diamonds$depth.in.mm))/(sd(diamonds$depth.in.mm))

lm_d_standard <- lm(price ~ carat_standardized + depth_standardized + table_standardized + length_standardized + width_standardized + depthmm_standardized)
tidy(lm_d_standard) %>%
  select(term, estimate, std.error)
```

A tibble: 7 × 3

	term	estimate	std.error
	<chr>	<dbl>	<dbl>
1	(Intercept)	3933.	6.45
2	carat_standardized	5065.	30.0
3	depth_standardized	-291.	7.89
4	table_standardized	-229.	6.89
5	length_standardized	-1476.	48.3
6	width_standardized	75.7	29.2
7	depthmm_standardized	29.4	31.3

You can not interpret standardized coefficients the same way as non standardized because the standardization transforms the values and now you can only see how each value is weighted in impact towards price.

Question 3


The seller should first look at the second model to determine how impactful each variable is in the model and see that carat size has the biggest impact on price. Then the seller should input the different carat sizes into the first model to get an estimated price to negotiate with the diamond store clerk. Only the first model can be used to estimate price.

Question 4

```
diamonds$cut = factor(diamonds$cut, levels = c("Fair", "Good", "Very Good", "Premium", "Ideal"))
diamonds$color = factor(diamonds$color, levels = c("J", "I", "H", "G", "F", "E", "D"))
diamonds$clarity = factor(diamonds$clarity, levels = c("I1", "SI1", "SI2", "VS1", "VS2", "VVS1",
```


Question 5

```
diamonds %>%
  ggplot(aes(x = carat, fill = color)) +
  geom_boxplot()
```



```
diamonds %>%
  ggplot(aes(x = price, fill = color)) +
```

```
geom_boxplot()
```



Question 6

```
lmcat <- lm(price ~ carat + color, data = diamonds)
tidy(lmcat)
```

A tibble: 8 × 5

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	-4051.	32.2	-126.	0
2	carat	8067.	14.0	575.	0
3	colorI	859.	34.3	25.1	8.13e-138
4	colorH	1182.	32.3	36.6	1.82e-289
5	colorG	1829.	31.5	58.0	0
6	colorF	1834.	32.2	57.0	0
7	colorE	1821.	32.3	56.4	0
8	colorD	1914.	33.8	56.7	0

price = intercept + 8066.62 * Carat + 858.74colorI + 1182.23colorH + 1828.93 colorG + 1834.20colorF + 1820.69 * colorE + 1914.47 * colorD

All the color variables are indicators that compare their value vs the baseline colorJ, so colorD = 1 if the Color of the Diamond is D and 0 otherwise. colorD increases the price of the Diamond by \$1914.47 compared to the price if the Diamond was colorJ.

Question 7

As the color increases/gets closure to D the price increases with the exception of ColorE where there is a slight dip, but this could be the result of a small sample size skewing the data because there is more higher carat colorF. This suggests that color has a postive association with price.

Question 8

```
lmint <- lm(price ~ carat:color, data = diamonds)
tidy(lmint)
```

```
# A tibble: 8 × 5
  term          estimate std.error statistic p.value
<chr>         <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept)   -2557.      12.9     -198.      0
2 carat:colorJ    6847.      22.8      300.      0
3 carat:colorI    7524.      19.3      390.      0
4 carat:colorH    7699.      18.6      414.      0
5 carat:colorG    8508.      19.9      427.      0
6 carat:colorF    8563.      22.4      383.      0
7 carat:colorE    8500.      24.6      346.      0
8 carat:colorD    8638.      28.0      308.      0
```

The interaction term between carat and colorD can be interpreted as a one unit increase in carat results in a 8638.10 increase in price if the color is colorD.


Question 9

```
lm_d_standard2 <- lm(price ~ length_standardized + width_standardized + depthmm_standardized, data = diamonds)
tidy(lm_d_standard2) %>%
  select(term, estimate, std.error)
```


```
# A tibble: 4 × 3
  term          estimate std.error
<chr>         <dbl>    <dbl>
1 (Intercept)   3933.      8.01
2 length_standardized 3130.     46.0
3 width_standardized  250.     36.1
4 depthmm_standardized  159.     33.6
```

```
diamonds%>%
  ggplot(aes(x = length.in.mm, y= price)) +
```

```
geom_point()
```



```
diamonds_filter <- diamonds %>% filter(length.in.mm > 0)
diamonds_filter %>%
  ggplot(aes(x = length.in.mm, y= price)) +
  geom_point()
```

```
lm_d2 <- lm(price ~ length.in.mm + width.of.mm, data = diamonds_filter)
tidy(lm_d2)
```

A tibble: 3 × 5

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	-14194.	41.3	-343.	0
2 length.in.mm	2958.	31.8	93.0	0
3 width.of.mm	204.	31.2	6.53	6.64e-11

Looking at the standardized model we can see that length.in.mm is contributing the most to the price by far with these three variables, with a value more than 10 times the next closest, so these other variables don't add much to the model, but I'll include width so there is a second variable. Then looking at the scatter plot we should filter by length > 0 to eliminate those outliers. The model is just price = intercept + 2958.266 * length.in.mm + 203.848 * width.of.mm

Question 10

```
lm_d2 <- lm(price ~ length.in.mm + width.of.mm, data = diamonds_filter)
tidy(lm_d2)
```

A tibble: 3 × 5

term	estimate	std.error	statistic	p.value
------	----------	-----------	-----------	---------

	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)		-14194.	41.3	-343.	0
2 length.in.mm		2958.	31.8	93.0	0
3 width.of.mm		204.	31.2	6.53	6.64e-11

A one unit increase in length.in.mm leads to a 2958.266 increase in price

Inference Wine

Question 1

```
wine <- read_csv("winequality-red.csv")
```

Rows: 1599 Columns: 12

— Column specification —

Delimiter: ","

dbl (12): fixed.acidity, volatile.acidity, citric.acid, residual.sugar, chlo...

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
lmW <- lm(quality ~ fixed.acidity + volatile.acidity + citric.acid + residual.sugar + chlorides +
tidy(lmW) %>%
  select(term, estimate, std.error)
```

A tibble: 12 × 3

	term	estimate	std.error
	<chr>	<dbl>	<dbl>
1	(Intercept)	22.0	21.2
2	fixed.acidity	0.0250	0.0259
3	volatile.acidity	-1.08	0.121
4	citric.acid	-0.183	0.147
5	residual.sugar	0.0163	0.0150
6	chlorides	-1.87	0.419
7	free.sulfur.dioxide	0.00436	0.00217
8	total.sulfur.dioxide	-0.00326	0.000729
9	density	-17.9	21.6
10	pH	-0.414	0.192
11	sulphates	0.916	0.114
12	alcohol	0.276	0.0265

Question 2

```
tidy(lmW)
```

A tibble: 12 × 5

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	22.0	21.2	1.04	3.00e- 1
2	fixed.acidity	0.0250	0.0259	0.963	3.36e- 1
3	volatile.acidity	-1.08	0.121	-8.95	9.87e-19
4	citric.acid	-0.183	0.147	-1.24	2.15e- 1
5	residual.sugar	0.0163	0.0150	1.09	2.76e- 1
6	chlorides	-1.87	0.419	-4.47	8.37e- 6
7	free.sulfur.dioxide	0.00436	0.00217	2.01	4.47e- 2
8	total.sulfur.dioxide	-0.00326	0.000729	-4.48	8.00e- 6
9	density	-17.9	21.6	-0.827	4.09e- 1
10	pH	-0.414	0.192	-2.16	3.10e- 2
11	sulphates	0.916	0.114	8.01	2.13e-15
12	alcohol	0.276	0.0265	10.4	1.12e-24

Null Hypothesis for pH would be $pH = 0$, the test stat is equal to the estimate/std.error, so $-0.4136/1.91e-01$ the p-value is 3.1×10^{-2} the t-stat is above a magnitude of 2 so meaning that we should reject the null. A t-stat of 2 is statistically significant to the 5% level.

Question 3

```
library(dplyr)

# Bootstrap function
MAKE_LM_BOOTSTRAP_STATS <- function(x_columns, y) {
  indices <- seq(from = 1, to = nrow(x_columns), by = 1)
  selected <- sample(x = indices, size = length(indices), replace = TRUE)

  boot_x <- x_columns %>%
    slice(selected)
  boot_y <- y[selected]

  boot <- data.frame(boot_x, y = boot_y)

  linear_model <- lm(y ~ ., data = boot)
  stat <- coef(linear_model)

  return(stat)
}

x <- wine %>% select(fixed.acidity , volatile.acidity , citric.acid , residual.sugar , chlorides)
y <- wine$quality

# Generate bootstrap distributions
bootstrap_distributions <-
  as.data.frame(replicate(n = 1000, MAKE_LM_BOOTSTRAP_STATS(x_columns = x, y = y)))

bootstrap_fixed <- as.numeric(bootstrap_distributions["fixed.acidity",])
bootstrap_volatile <- as.numeric(bootstrap_distributions["volatile.acidity",])
```

```
bootstrap_citric <- as.numeric(bootstrap_distributions["citric.acid",])
bootstrap_sugar <- as.numeric(bootstrap_distributions["residual.sugar",])
bootstrap_chlorides <- as.numeric(bootstrap_distributions["chlorides",])
bootstrap_free <- as.numeric(bootstrap_distributions["free.sulfur.dioxide",])
bootstrap_total <- as.numeric(bootstrap_distributions["total.sulfur.dioxide",])
bootstrap_density <- as.numeric(bootstrap_distributions["density",])
bootstrap_pH <- as.numeric(bootstrap_distributions["pH",])
bootstrap_sulphates <- as.numeric(bootstrap_distributions["sulphates",])
bootstrap_alcohol <- as.numeric(bootstrap_distributions["alcohol",])
```

```
# Compute 95% confidence intervals
```

```
CI_fixed <- quantile(bootstrap_fixed, c(0.025, 0.975))
CI_volatile <- quantile(bootstrap_volatile, c(0.025, 0.975))
CI_citric <- quantile(bootstrap_citric, c(0.025, 0.975))
CI_chlorides <- quantile(bootstrap_chlorides, c(0.025, 0.975))
CI_sugar <- quantile(bootstrap_sugar, c(0.025, 0.975))
CI_free <- quantile(bootstrap_free, c(0.025, 0.975))
CI_total <- quantile(bootstrap_total, c(0.025, 0.975))
CI_density <- quantile(bootstrap_density, c(0.025, 0.975))
CI_pH <- quantile(bootstrap_pH, c(0.025, 0.975))
CI_sulphates <- quantile(bootstrap_sulphates, c(0.025, 0.975))
CI_alcohol <- quantile(bootstrap_alcohol, c(0.025, 0.975))
print(CI_fixed)
```

```
      2.5%      97.5%
-0.03791511  0.08506279
```

```
print(CI_volatile)
```

```
      2.5%      97.5%
-1.3399732 -0.8070485
```

```
print(CI_citric)
```

```
      2.5%      97.5%
-0.4590801  0.1030521
```

```
print(CI_chlorides)
```

```
      2.5%      97.5%
-2.8197102 -0.9443352
```

```
print(CI_sugar)
```

```
      2.5%      97.5%
-0.02250166  0.05643550
```

```
print(CI_free)
```

```
      2.5%      97.5%
-3.745692e-05  8.602863e-03
```

```
print(CI_total)
```

```
      2.5%      97.5%
-0.004773532 -0.001926829
```

```
print(CI_density)
```

```
      2.5%      97.5%
-67.25123   32.00337
```

```
print(CI_pH)
```

```
      2.5%      97.5%
-0.8236101273  0.0008394741
```

```
print(CI_sulphates)
```

```
      2.5%      97.5%
0.6815787  1.2214198
```

```
print(CI_alcohol)
```

```
      2.5%      97.5%
0.2216677  0.3303539
```

```
confint(lmW, level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	-1.960710e+01	63.537517843
fixed.acidity	-2.590639e-02	0.075887499
volatile.acidity	-1.321126e+00	-0.846054953
citric.acid	-4.712441e-01	0.106116245
residual.sugar	-1.309474e-02	0.045757280
chlorides	-2.696632e+00	-1.051817956
free.sulfur.dioxide	1.024314e-04	0.008620235
total.sulfur.dioxide	-4.693951e-03	-0.001835208
density	-6.031362e+01	24.551294542
pH	-7.894637e-01	-0.037842600
sulphates	6.920661e-01	1.140602768
alcohol	2.242512e-01	0.328144192

Question 4

Sulphates, free sulfur and Alcohol have a positive effect on quality while, sugar, fixed.acidity and density can have a positive or negative effect depending on what end of the CI they are on. The results between the bootstrap and the parametric are fairly similar with the bounds being slightly different, but the the variables with a positive effect are all the same in both Confidence Intervals.

Question 5

```
lm1 <- lm(quality ~ fixed.acidity + volatile.acidity + citric.acid + chlorides + free.sulfur.diox:
lm1 %>%
  glance() %>%
  select(df, df.residual, statistic, p.value)
```

A tibble: 1 × 4

	df	df.residual	statistic	p.value
	<dbl>	<int>	<dbl>	<dbl>
1	9	1589	99.3	3.31e-147

```
lmW %>%
  glance() %>%
  select(df, df.residual, statistic, p.value)
```

A tibble: 1 × 4

	df	df.residual	statistic	p.value
	<dbl>	<int>	<dbl>	<dbl>
1	11	1587	81.3	1.79e-145

Ho the variables residual.sugar and density are 0 and have no effect on quality

The model with residual.sugar and density has 2 more degrees of freedom and 2 more residuals, but is less accurate with a lower p.value and t-statistic. Therefore we should accept the null and not include these variables in the regression model.


Diagnostics

Question 1

```
diamonds$residuals <- residuals(lmint)
diamonds$fitted_values <- fitted(lmint)


ggplot(diamonds, aes(x = price)) +
  geom_density(fill = "blue", alpha = 0.5) +
  labs(x = "Price", y = "Density", title = "Density Plot of Diamond Prices")
```

Density Plot of Diamond Prices



```
ggplot(diamonds, aes(x = fitted_values, y = residuals)) +  
  geom_point(alpha = 0.5) +  
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +  
  labs(x = "Fitted Values", y = "Residuals", title = "Residuals vs Fitted Plot")
```

Residuals vs Fitted Plot




```
residuals_standardized <- rstandard(lmint)

qq_data <- qqnorm(residuals_standardized, plot.it = FALSE)

qq_df <- data.frame(Theoretical = qq_data$x, StandardizedResiduals = qq_data$y)

ggplot(qq_df, aes(x = Theoretical, y = StandardizedResiduals)) +
  geom_point() +
  geom_abline(slope = 1, intercept = 0, color = "blue") +
  labs(title = "QQ Plot of Standardized Residuals",
       x = "Theoretical Quantiles",
       y = "Standardized Residuals") +
  theme_minimal()
```


QQ Plot of Standardized Residuals



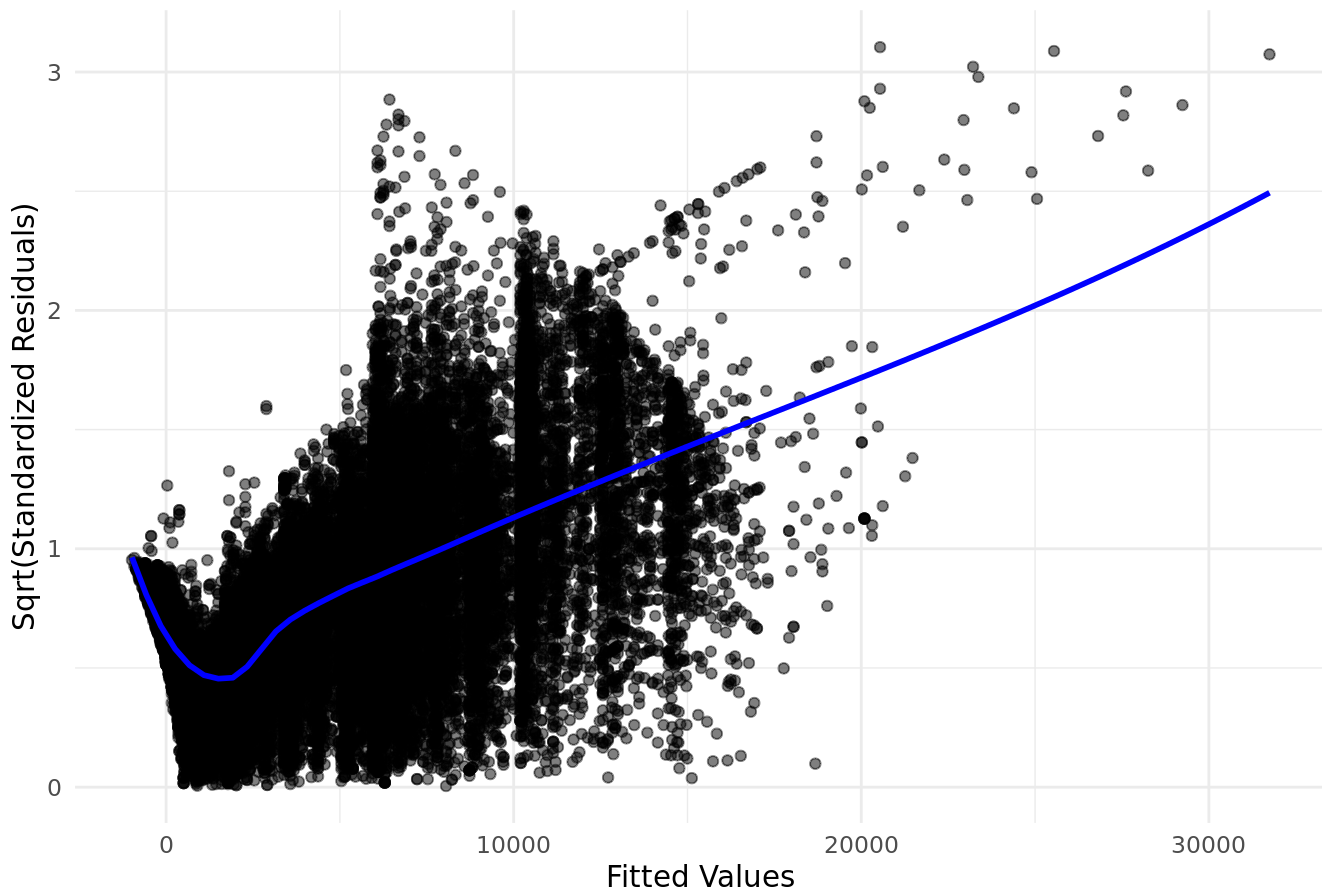
```
sqrt_std_residuals <- sqrt(abs(rstandard(lmint)))

scale_location_df <- data.frame(
  FittedValues = fitted(lmint),
  SqrtStdResiduals = sqrt_std_residuals
)

ggplot(scale_location_df, aes(x = FittedValues, y = SqrtStdResiduals)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess", color = "blue", se = FALSE) +
  labs(title = "Scale-Location Plot",
       x = "Fitted Values",
       y = "Sqrt(Standardized Residuals)") +
  theme_minimal()
```

`geom_smooth()` using formula = 'y ~ x'

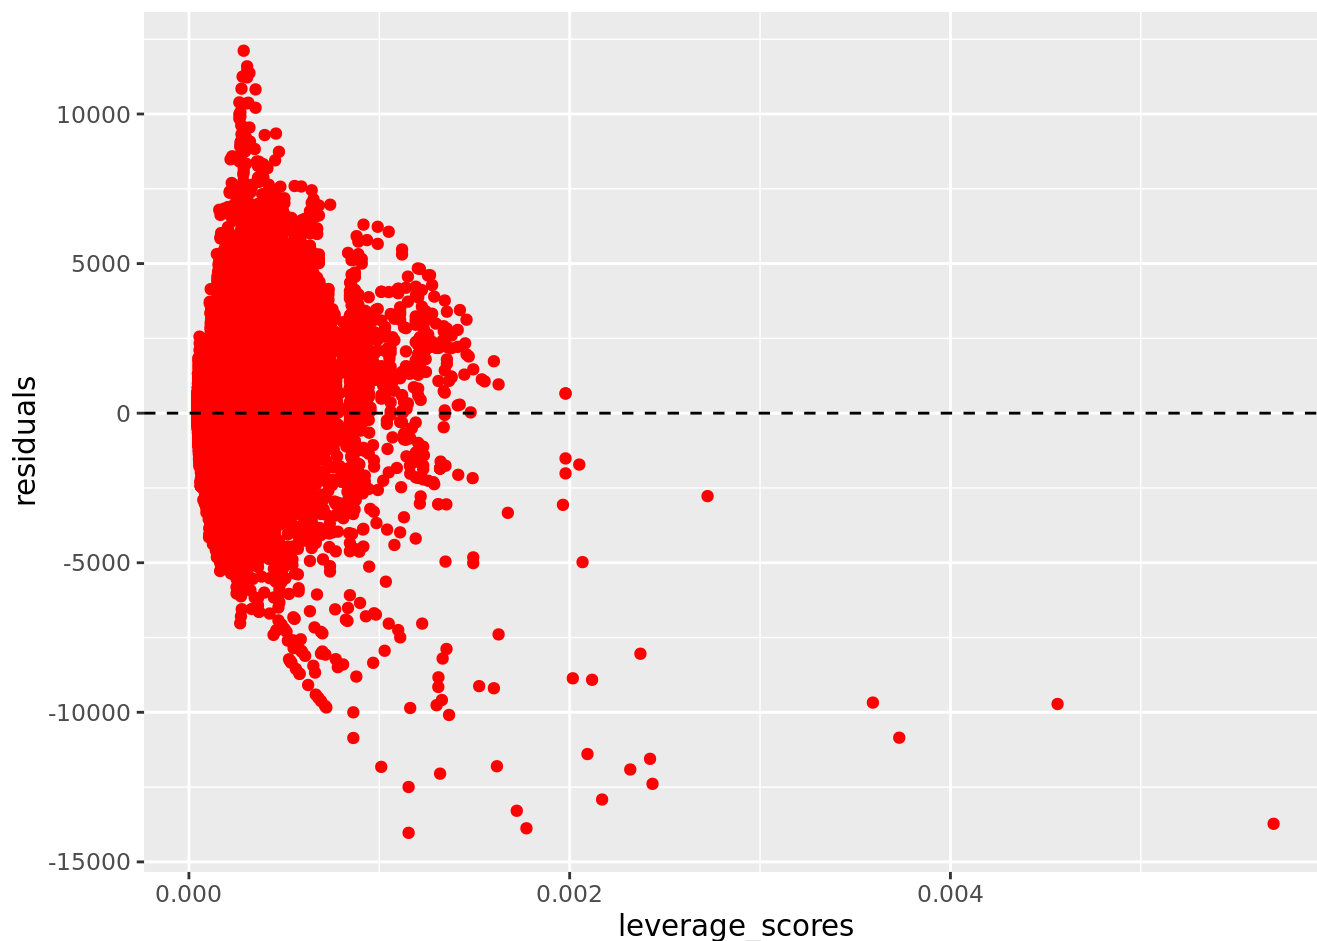
Scale-Location Plot



These three graphs show that this diamonds data set violates most assumptions, the residuals vs fitted graph shows that there is a ton of variance that isn't evenly distributed especially past 15000 violating homoscedasticity assumption. The QQ plot shows that the error terms are not normally distributed around the tails of the graph swinging wildly in an S like pattern. The scale location plot shows again the problem with uneven variance.

Question 2

```
leverage_scores <- hatvalues(lmint)
diamonds %>%
  ggplot(aes(x = leverage_scores, y = residuals)) +
  geom_point(color = "red") +
  geom_hline(yintercept = 0, linetype = "dashed")
```



```
diamonds %>%
  mutate(leverages = hatvalues(lmint)) %>%
  filter(leverages > 0.002)
```

A tibble: 15 × 19


	carat	cut	color	clarity	depth	table	price	length.in.mm	width.of.mm
	<dbl>	<fct>	<fct>	<fct>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	3.11	Fair	J	I1	65.9	57	9823	9.15	9.02
2	3.05	Premium	E	I1	60.9	58	10453	9.26	9.25
3	3	Good	E	I1	64.2	65	11548	9.08	8.96
4	2.75	Ideal	D	I1	60.9	57	13156	9.04	8.98
5	2.58	Very Good	D	SI2	58.9	63	14749	9.08	9.01
6	4.01	Premium	I	I1	61	61	15223	10.1	10.1
7	4.01	Premium	J	I1	62.5	62	15223	10.0	9.94
8	3.4	Fair	D	I1	66.8	52	15964	9.42	9.34
9	4	Very Good	I	I1	63.3	58	15984	10.0	9.94
10	3.67	Premium	I	I1	62.4	56	16193	9.86	9.81
11	4.13	Fair	H	I1	64.8	61	17329	10	9.85
12	2.57	Premium	D	SI2	58.9	58	17924	8.99	8.94
13	5.01	Fair	J	I1	65.5	59	18018	10.7	10.5
14	4.5	Fair	J	I1	65.8	58	18531	10.2	10.2
15	3.51	Premium	J	VS2	62.5	59	18701	9.66	9.63

i 10 more variables: depth.in.mm <dbl>, carat_standardized <dbl>,

```
# depth_standardized <dbl>, table_standardized <dbl>,  
# length_standardized <dbl>, width_standardized <dbl>,  
# depthmm_standardized <dbl>, residuals <dbl>, fitted_values <dbl>,  
# leverages <dbl>
```

```
cooks_distance <- cooks.distance(lmint)  
  
plot(cooks_distance, type="h", main="Cook's Distance", ylab="Cook's Distance")  
abline(h=4/(nrow(diamonds)), col="red")
```

Cook's Distance



```
threshold <- 4 / length(cooks_distance)  
  
outliers <- cooks_distance > threshold  
  
diamonds_cleaned <- diamonds[!outliers, ]
```

Based on the graph we will say that points above 0.002 have a high leverage, filtering by this we can see that diamonds with fair cut and J color have the highest leverage, so there could be something different about these diamonds that is causing a lot of error. A diamond with a fair cut and J color should typically be a lower price so something about these diamonds makes them much more expensive than a normal J color fair cut diamond.

Prediction

Question 1

```
d3 <- diamonds %>% mutate(ln_width = log(width.of.mm),
                          ln_price = log(price))
sum(is.infinite(d3$lnwidth))
```

Warning: Unknown or uninitialised column: `lnwidth`.

```
[1] 0
```

```
sum(is.infinite(d3$lnprice))
```

Warning: Unknown or uninitialised column: `lnprice`.

```
[1] 0
```

```
sum(is.nan(d3$ln_width))
```

```
[1] 0
```

```
sum(is.nan(d3$ln_price))
```

```
[1] 0
```

```
sum(is.na(d3$ln_width))
```

```
[1] 0
```

```
sum(is.na(d3$ln_price))
```

```
[1] 0
```

```
lm3 <- lm(log(price) ~ width.of.mm, data = diamonds)
tidy(lm3)
```

A tibble: 2 × 5

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	3.02	0.00786	384.	0
2 width.of.mm	0.832	0.00134	618.	0

I tried to take add the ln(price) and ln(width) to the diamonds data set then make a linear model of it but it kept giving me an error saying that Error in lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :

NA/NaN/Inf in 'x', so I'm just gonna interpret the regular regression and pretend its log log instead of log lin. For every 1% increase in width.of.mm there is a 0.83% increase in price holding all other variables constant.

Question 2

```
fake_data <- data.frame(width.of.mm = c(7,20))

predict(object = lm3, newdata = fake_data)
```

```
      1      2
8.839234 19.651036
```

```
predict(object = lm3, newdata = fake_data, interval = "prediction")
```

```
      fit      lwr      upr
1 8.839234 8.140099 9.53837
2 19.651036 18.950898 20.35117
```

Based on these intervals we can say with 95% confidence that the predicted price is between 8.14 and 9.53 for the first observation and 18.95 and 20.35 for the second.

Question 3

```
lm3 %>%
  glance() %>%
  select(r.squared)
```

```
# A tibble: 1 × 1
  r.squared
  <dbl>
1    0.876
```

The bounds for both prediction intervals are both about 1.4 so its hard to say, that one is more untrustworthy, we could fit an r-squared graph to see how accurate the prediction is when you add more variables, but we only have one variable and the starting r.squared it pretty high with 0.876 suggesting that most of the variability in price comes from width.of.mm.

Wine

Question 1

```
set.seed(1234)
train_pct <- .7
indices <- seq(from = 1, to = nrow(wine), by = 1)
```

```

training_indices <- sample(x = indices, replace = FALSE,
                           size = nrow(wine)*train_pct)

wine_train <- wine %>%
  slice(training_indices)
wine_test <- wine %>%
  slice(-training_indices)

lmW1 <- lm(quality ~ alcohol, data = wine_train)
lmW2 <- lm(quality ~ alcohol + sulphates, data = wine_train)
lmW3 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide , data = wine_train)
lmW4 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity , data = wine_train)
lmW5 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity + volatile.acidity
lmW6 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity + volatile.acidity
lmW7 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity + volatile.acidity
lmW8 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity + volatile.acidity
lmW9 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity + volatile.acidity
lmW10 <- lm(quality ~ alcohol + sulphates + free.sulfur.dioxide + fixed.acidity + volatile.acidity

predictions_1 <- predict(object = lmW1, newdata = wine_test)
predictions_2 <- predict(object = lmW2, newdata = wine_test)
predictions_3 <- predict(object = lmW3, newdata = wine_test)
predictions_4 <- predict(object = lmW4, newdata = wine_test)
predictions_5 <- predict(object = lmW5, newdata = wine_test)
predictions_6 <- predict(object = lmW6, newdata = wine_test)
predictions_7 <- predict(object = lmW7, newdata = wine_test)
predictions_8 <- predict(object = lmW8, newdata = wine_test)
predictions_9 <- predict(object = lmW9, newdata = wine_test)
predictions_10 <- predict(object = lmW10, newdata = wine_test)

best_k_model <- regsubsets(quality ~ ., wine)
summary(best_k_model)

```

Subset selection object

Call: regsubsets.formula(quality ~ ., wine)

11 Variables (and intercept)

	Forced in	Forced out
fixed.acidity	FALSE	FALSE
volatile.acidity	FALSE	FALSE
citric.acid	FALSE	FALSE
residual.sugar	FALSE	FALSE
chlorides	FALSE	FALSE
free.sulfur.dioxide	FALSE	FALSE
total.sulfur.dioxide	FALSE	FALSE
density	FALSE	FALSE
pH	FALSE	FALSE
sulphates	FALSE	FALSE
alcohol	FALSE	FALSE

1 subsets of each size up to 8

Selection Algorithm: exhaustive

fixed.acidity volatile.acidity citric.acid residual.sugar chlorides

1	(1)	" "	" "	" "	" "	" "
2	(1)	" "	"*"	" "	" "	" "
3	(1)	" "	"*"	" "	" "	" "
4	(1)	" "	"*"	" "	" "	" "
5	(1)	" "	"*"	" "	" "	"*"
6	(1)	" "	"*"	" "	" "	"*"
7	(1)	" "	"*"	" "	" "	"*"
8	(1)	" "	"*"	"*"	" "	"*"
free.sulfur.dioxide total.sulfur.dioxide density pH sulphates alcohol						
1	(1)	" "	" "	" "	" " " "	"*"
2	(1)	" "	" "	" "	" " " "	"*"
3	(1)	" "	" "	" "	" " "*"	"*"
4	(1)	" "	"*"	" "	" " "*"	"*"
5	(1)	" "	"*"	" "	" " "*"	"*"
6	(1)	" "	"*"	" "	"* " "*"	"*"
7	(1)	"*"	"*"	" "	"* " "*"	"*"
8	(1)	"*"	"*"	" "	"* " "*"	"*"

Question 2

```

k <- 10
fold_vector <-
  cut(1:nrow(wine), breaks = k, labels = FALSE)

random_folds <- sample(x = fold_vector, size = nrow(wine),
                      replace = FALSE)

wine <- wine %>%
  mutate(folds = random_folds)
ONE_CV_FOLD <- function(fold_number, formula_string)
{
  wine_train <- wine %>%
    filter(folds != fold_number)

  wine_test <- wine %>%
    filter(folds == fold_number)

  cv_linear_model <-
    do.call(what = "lm",
            args = list(formula = as.formula(formula_string), data = quote(wine_train)))

  cv_predictions <- predict(object= cv_linear_model, newdata = wine_test)

  observations <- wine_test %>%
    select(quality) %>%
    pull()

  MSE <- mean((cv_predictions-observations)^2)

  return(MSE)

```



```
}  
MSEs1 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol")  
mean(MSEs1)
```

[1] 0.5057258

```
MSEs2 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity")  
mean(MSEs2)
```

[1] 0.4470504

```
MSEs3 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates")  
mean(MSEs3)
```

[1] 0.4351985

```
MSEs4 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates + total.sulfur.d:  
mean(MSEs4)
```

[1] 0.4308383

```
MSEs5 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates + total.sulfur.d:  
mean(MSEs5)
```

[1] 0.4265202

```
MSEs6 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates + pH + chlorides  
mean(MSEs6)
```

[1] 0.4231016

```
MSEs7 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates + total.sulfur.d:  
mean(MSEs7)
```

[1] 0.4231016

```
MSEs8 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates + total.sulfur.d:  
mean(MSEs8)
```

[1] 0.4222518

```
MSEs9 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                formula_string = "quality~ alcohol + volatile.acidity + sulphates + total.sulfur.d:  
mean(MSEs9)
```

[1] 0.4222533

```
MSEs10 <- sapply(unique(fold_vector), FUN = ONE_CV_FOLD,  
                 formula_string = "quality~ alcohol + volatile.acidity + sulphates + total.sulfur.d:  
mean(MSEs10)
```

[1] 0.422882

Based on the MSEs decreasing with every variable added I would say the first model with only alcohol is the "best" model